

**ECE4893A/CS4803MPG:  
MULTICORE AND GPU  
PROGRAMMING  
FOR VIDEO GAMES**

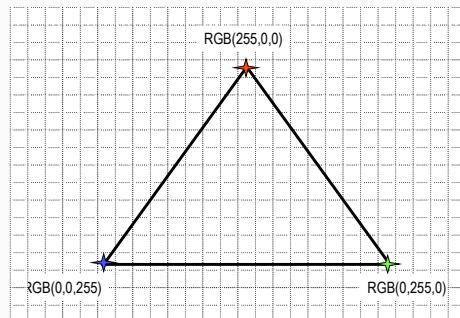
3D Rendering Pipeline (III)




Prof. Hsien-Hsin Sean Lee  
 School of Electrical and Computer  
 Engineering  
 Georgia Institute of Technology



## Rasterization: Shading a Triangle

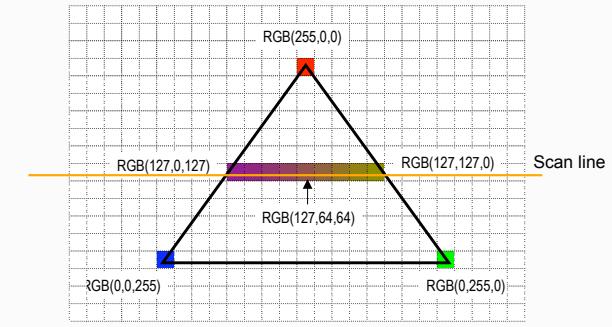


- Converting geometry to a raster image (i.e., pixels)
- Paint each pixel's color (by calculating light intensity) on your display (Typically object-based)
- Gouraud Shading: Intensity Interpolation of vertices

Georgia Institute  
of Technology

2

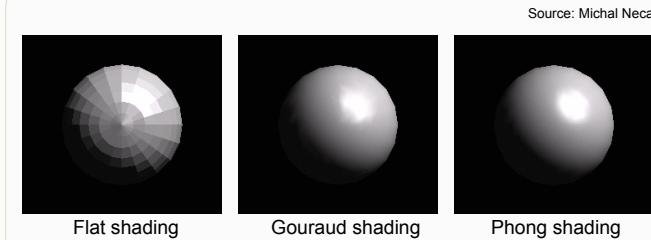
**Gouraud Shading**



- Scan conversion algorithm



## Comparison of Shading Methods

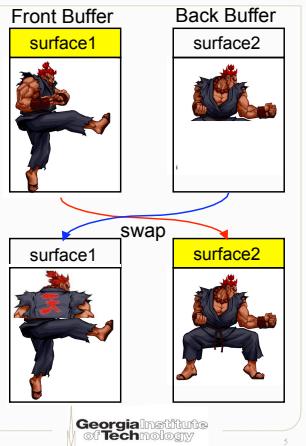


- Phong shading
  - requires generating per-pixel normals to compute light intensity for each pixel, not efficient for games
  - Can be done on GPGPU today using Cg or HLSL
- Gouraud shading is supported by Graphics hardware

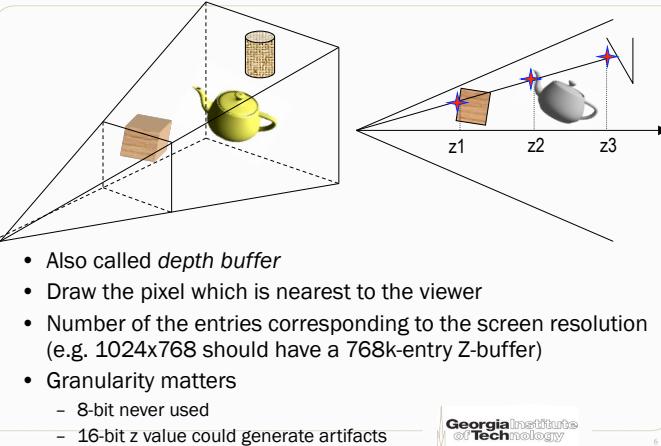
Georgia Institute  
of Technology

## Double Buffering

- Display refreshes at 60 ~ 75 Hz
- Rendering could be “faster” than the refresh period
- Too fast leads to
  - Frames not shown
- Too slow leads to
  - New and old frame mixed
  - Flickering
- Solution:
  - *Double or multiple buffering*



## Z-Buffer

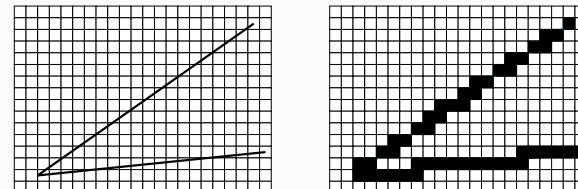


- Also called *depth buffer*
- Draw the pixel which is nearest to the viewer
- Number of the entries corresponding to the screen resolution (e.g. 1024x768 should have a 768k-entry Z-buffer)
- Granularity matters
  - 8-bit never used
  - 16-bit z value could generate artifacts

## Z-Buffer Bandwidth Could Be an Issue

- Perform Z test before drawing a pixel
- How much bandwidth needed for  $RZ + WZ + RC + WC + TR$  per pixel?
- Overdrawn rate is about 3 out of 4
  - One every 4 pixels drawn is for clearing the Z-buffer
- Some cards perform Z-compression to alleviate bandwidth issues

## Aliasing



- Jagged line (or staircase)
- Can be improved by increasing resolution (i.e. more pixels)

## Anti-Aliasing by Multisampling (Example: Supersampling)

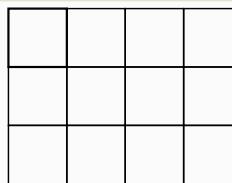
1	2	1
2	4	2
1	2	1

3x3 Virtual Pixels  
(Bartlett window)

(255,255,255) (255,255,255) (255,0,0)  
 (255,255,255) (255,0,0) (255,255,255)  
 (255,0,0) (255,255,255) (255,255,255)

Example

(255, 159, 159)



Actual Screen Pixels

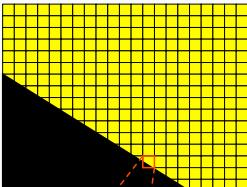
- GPU samples multiple locations for a pixel
- Several different methods
  - e.g., grid (as shown), random, GeForce's quincunx
- Downside
  - Blurry image
  - Increased memory (e.g., z-buffer) storage for subpixel information

Georgia Institute of Technology

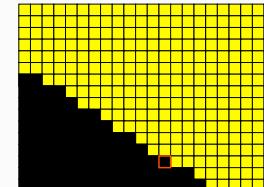
9

## Anti-Aliasing Example

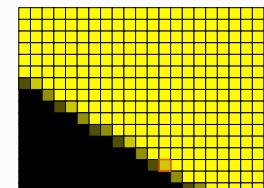
Ideal



No MSAA



With MSAA

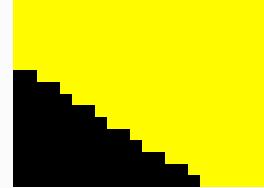


Georgia Institute of Technology

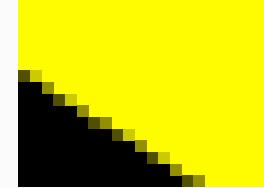
10

## Visualizing Anti-Aliasing Example

No MSAA



With MSAA



Georgia Institute of Technology

11

## Texture Mapping

- Rendering tiny triangles is slow
- Players won't even look at some certain details
  - Sky, clouds, walls, terrain, wood patterns, etc.
- Simple way to add details and enhance realism
- Use 2D images to map polygons
- Images are composed of 2D "texels"
- Can be used to substitute or blend the lit color of a texture-mapped surface

Georgia Institute of Technology

12

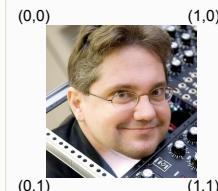
## Texture Mapping

- Introduce one more component to geometry
  - Position coordinates
  - Normal vector
  - Color
  - *Texture coordinates*
- Texture info
  - $(u, v)$  coordinates for each vertex
  - Typically range from 0 to 1
    - (0,0) from upper left corner



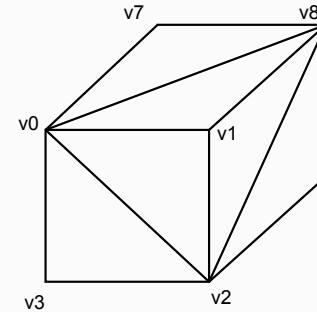
13

## Texture Mapping



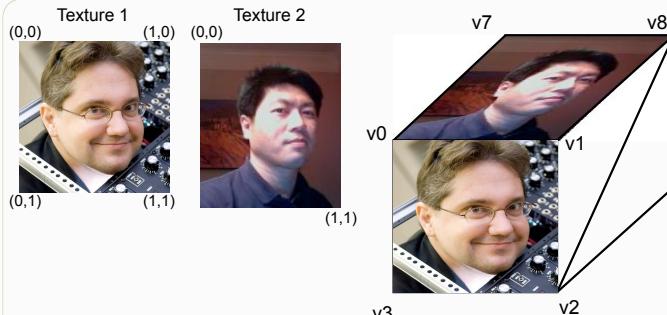
Texture: hunk.jpg

$\downarrow$        $\downarrow$   
 $\{v1.x, v1.y, v1.z, \dots, 1, 0\},$   
 $\{v2.x, v2.y, v2.z, \dots, 1, 1\},$   
 $\{v0.x, v0.y, v0.z, \dots, 0, 0\},$   
 $\{v3.x, v3.y, v3.z, \dots, 0, 1\},$



14

## Texture Mapping

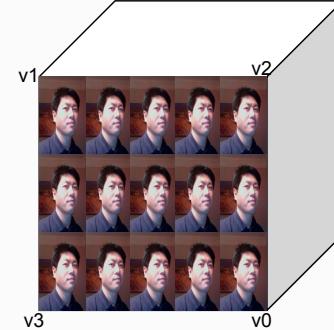


15

## Texture Mapping

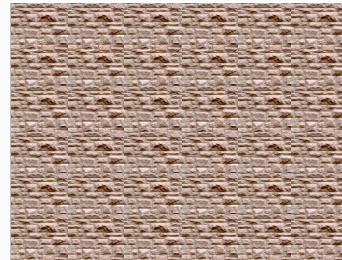


$\downarrow$        $\downarrow$   
 $\{v1.x, v1.y, v1.z, \dots, 0, 0\},$   
 $\{v2.x, v2.y, v2.z, \dots, 5, 0\},$   
 $\{v0.x, v0.y, v0.z, \dots, 5, 3\},$   
 $\{v3.x, v3.y, v3.z, \dots, 0, 3\},$



16

## Texture Mapping



u

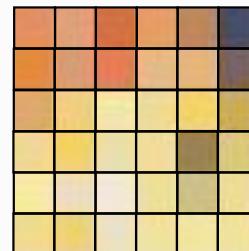
v

```
{v1.x, v1.y, v1.z, ..., 0, 0},
{v2.x, v2.y, v2.z, ..., 6, 0},
{v0.x, v0.y, v0.z, ..., 6, 6},
{v3.x, v3.y, v3.z, ..., 0, 6},
```

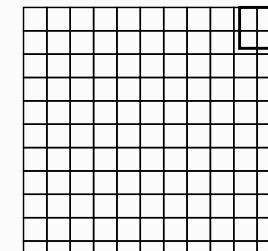
Georgia Institute  
of Tech nology

17

## Magnification



Texels



Pixels on screen

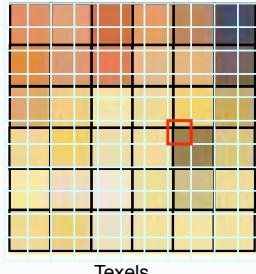
- Texel and pixel mapping is rarely 1-to-1
- Mapped triangle is very close to the camera
- One texel maps to multiple pixels



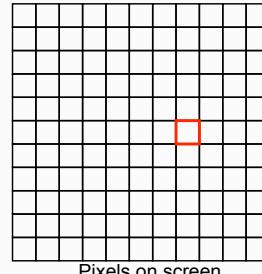
Georgia Institute  
of Tech nology

18

## Nearest Point Sampling (for Magnification)



Texels



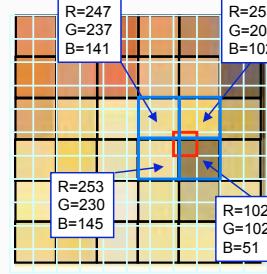
Pixels on screen

- Choose the texel nearest the pixel's center

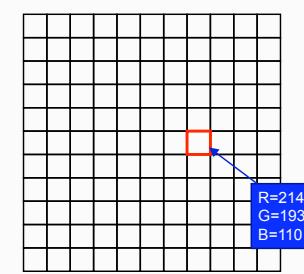
Georgia Institute  
of Tech nology

19

## Bi-linear Filtering (for Magnification)



Texels



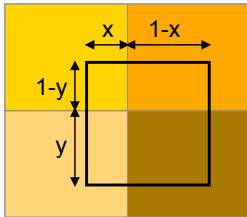
Pixels on screen

- Average for the 2x2 texels surrounding a given pixel

Georgia Institute  
of Tech nology

20

## Bi-linear Filtering (for Magnification)



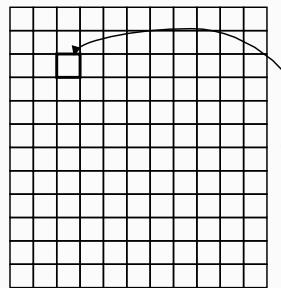
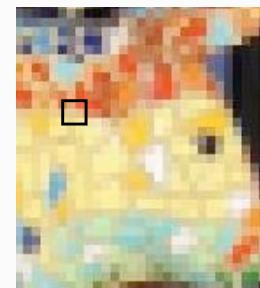
: pixel enclosed by 4 texels

- Or take the weighted color values for the 2x2 texels surrounding a given pixel



21

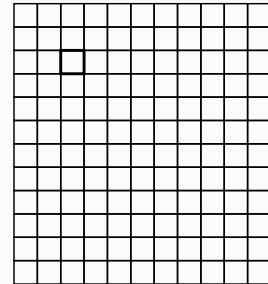
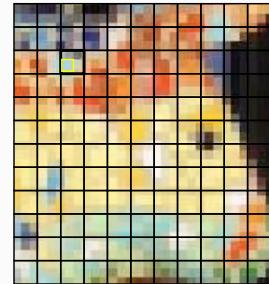
## Minification



22

- Texel and pixel mapping is rarely 1-to-1
- Multiple texels map to one pixel

## Nearest Point Sampling (for Minification)

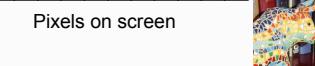
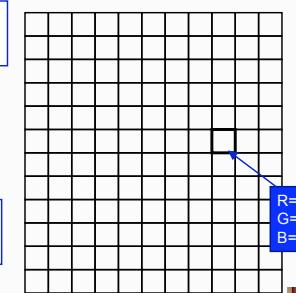
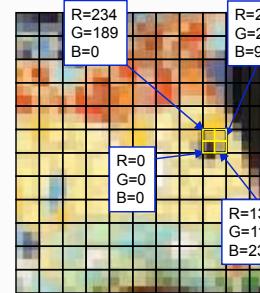


- Choose the texel nearest the pixel's center



23

## Bi-linear Filtering (for Minification)



24

- Average for the 2x2 texels surrounding a given pixel



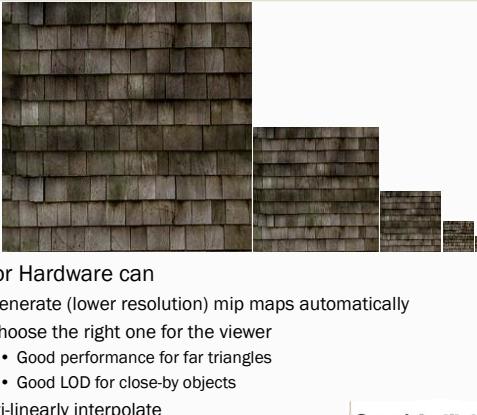
## Mip-mapping

- Multiple versions are provided for the same texture
- Different versions have different levels of details
  - E.g., 7 LOD maps: 256x256, 128x128, 64x64, 32x32, 16x16, 8x8, 4x4
  - Choose the closest maps to render a surface
- Maps can be automatically generated by 3D API
- Accelerate texture mapping for far-away polygons
- More space to store texture maps



25

## Mip-mapping

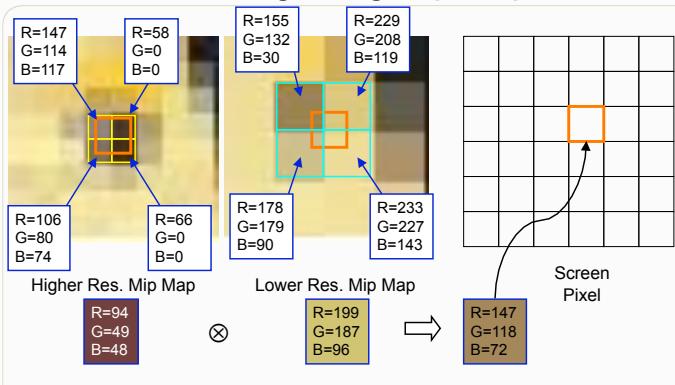


- API or Hardware can
  - Generate (lower resolution) mipmap automatically
  - Choose the right one for the viewer
    - Good performance for far triangles
    - Good LOD for close-by objects
  - Tri-linearly interpolate



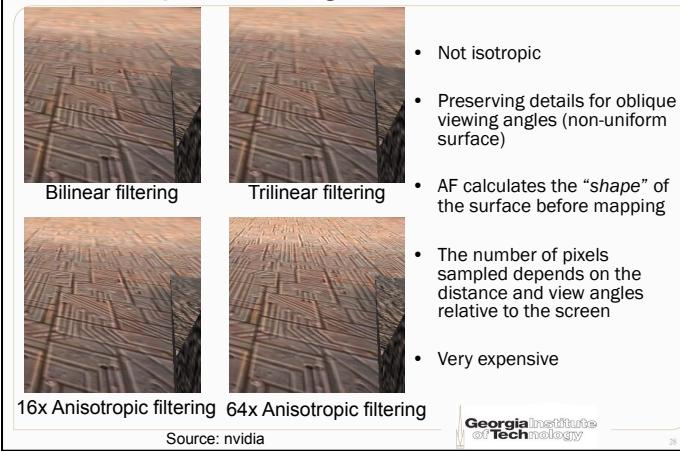
26

## Tri-linear Filtering using Mip maps



27

## Anisotropic Filtering



28

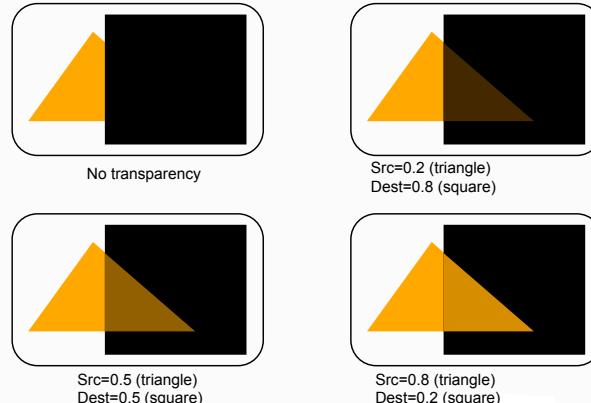
## Color Blending and Alpha Blending

- Transparency effect (e.g. Water, glasses, etc.)
- Source color blended with destination color
- Several blending methods
  - Additive  
 $C = \text{SrcPixel} \otimes (1,1,1,1) + \text{DstPixel} \otimes (1,1,1,1) = \text{SrcPixel} + \text{DstPixel}$
  - Subtractive  
 $C = \text{SrcPixel} \otimes (1,1,1,1) - \text{DstPixel} \otimes (1,1,1,1) = \text{SrcPixel} - \text{DstPixel}$
  - Multiplicative  
 $C = \text{DstPixel} \otimes \text{SrcPixel}$
  - Using Alpha value in the color (Alpha Blending)  
 $C = \text{SrcPixel} \otimes (\alpha,\alpha,\alpha,\alpha) + \text{DstPixel} \otimes (1-\alpha,1-\alpha,1-\alpha,1-\alpha)$
  - And many more in the API ...



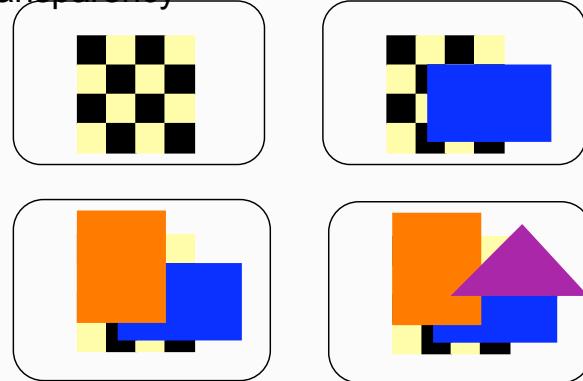
29

## Alpha Blending (Inverse Source form)



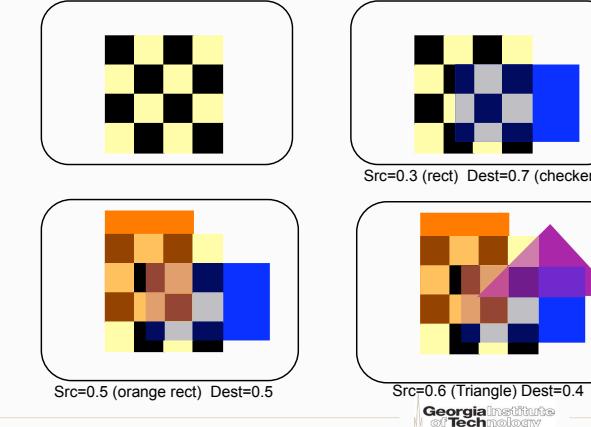
30

## Another Example Without Transparency



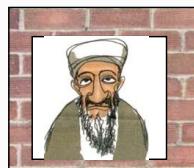
31

## Another Alpha Blending Example

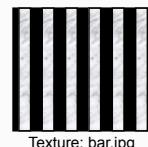


32

## Alpha Test



Straightforward texture mapping



- Reject pixels by checking their alpha values
- Model fences, chicken wires, etc.

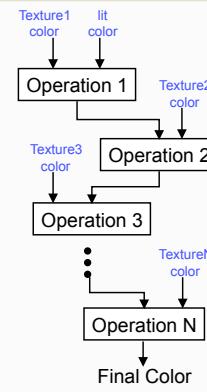
Georgia Institute of Technology

33

```
if (α < val)
    reject pixel
else
    accept pixel
```

## Multi-Texturing

- Map multiple texture to a polygon
  - Common APIs support 8 textures
- Performance will be reduced
- Multiple texturing stages in the pipeline
- Texture color will be calculated by
  - Multiplication
  - Addition
  - Subtraction



Georgia Institute of Technology

34

## Multi-Texturing Example: Light Mapping

 $\otimes$ 

Some crumpled paper texture

A spotlight map



Different alpha blending

Georgia Institute of Technology

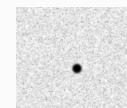
35

## Bump Mapping

- Per-fragment lighting using bump map (normal map) to perturb surface normal
- No geometry tessellation, avoid geometry complexity
- Store normal vectors rather than RGB colors for bump map
- Apply per-pixel shading (w/ light vector, e.g., Phong Shading)



Shaded sphere



Bump map



Bump mapped sphere

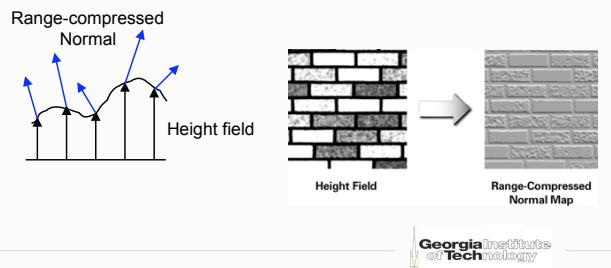
Source: wikipedia

Georgia Institute of Technology

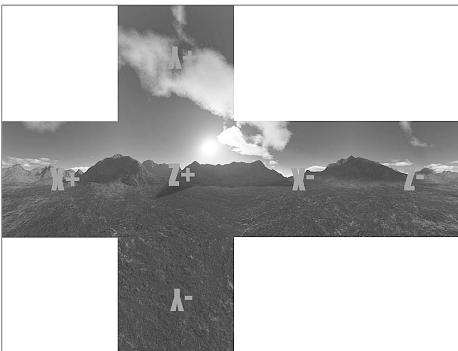
36

## Bump Mapping

- Normal map was derived from a Height field map
  - Height field stores the “elevation” for each texel
  - Sample texel’s height as well as texels to the right and above



## (Cube) Environment Mapping



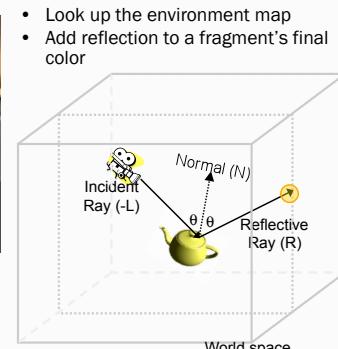
- Cube Map Textures (in World coordinate)
- Each face encodes 1/6 of the panoramic environment

Source: Game Creators, Ltd.

Georgia Institute of Technology

38

## Reflective (or Environment) Mapping



Function `texCUBE()` is provided for sampling the texel in cube map

$$R = 2 * (\text{dotp}(N, L)) * N - L$$

Georgia Institute of Technology

39

## Reflective Mapping



Cube texture map



Rendered image

Georgia Institute of Technology

40

## Stenciling

- Stencil buffer

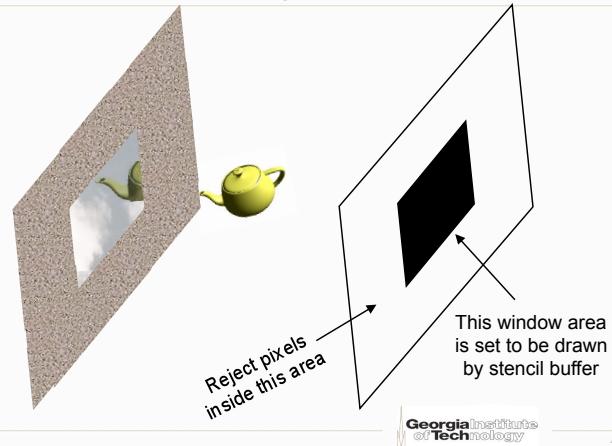
- To reject certain pixels to be displayed
- To create special effect similar to alpha test
  - Mask out part of the screen
- Set together with Z-buffer in 3D API
- Perform prior to Z-buffer test

```
if ((stencil ref & mask)
    op (pixel val & mask))
accept pixel
else
reject pixel
```

Georgia Institute  
of Tech nology

41

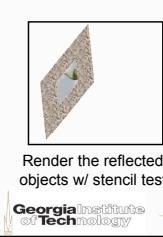
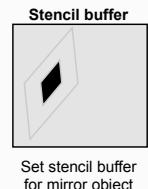
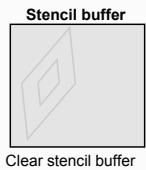
## Stencil Buffer Example



42

## Mirror Effect

1. Render the entire scene as normal (no reflection yet)
2. Clear the entire stencil buffer to '0' (i.e., mirror's fragments)
3. Render the mirror primitives and set the corresponding stencil buffer fragment to '1'
4. Render the reflected objects only if stencil test passes (i.e., value==1)
  - Using a "reflection matrix" for world transformation (Draw the scene as if they are seen in the mirror)



Georgia Institute  
of Tech nology

43

## Mirror Effect (Cont.)

Can be done in a reverse order

1. Render the reflected image of the scene using a "reflection matrix" for world transformation (Draw the scene as if they are seen in the mirror)
2. Render non-reflected with stencil buffer accept/reject test to prevent the reflected image being drawn over

Georgia Institute  
of Tech nology

44

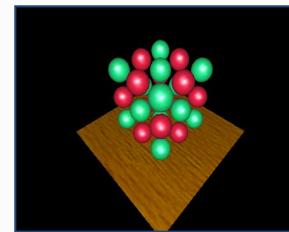
## Cast Shadow

- Light source is the eye
- Objects are not visible from the light are in shadow (but illuminated by ambient or emissive lights)
- Two-pass process proposed by Lance Williams in SIGGRAPH'98

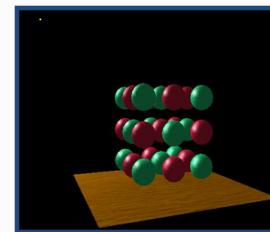
Georgia Institute  
of Technology

45

## Shadow Map



Scene from light's Point-of-view



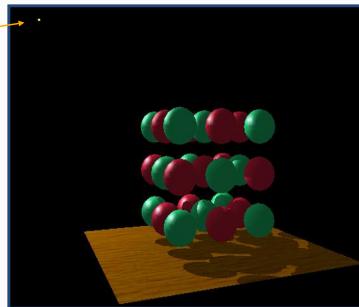
Scene from eye's Point-of-view

Georgia Institute  
of Technology

46

## Shadow Map

Light source



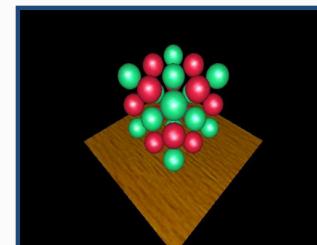
Desired rendered scene with shadow

Figure Source: Justin Reynolds

Georgia Institute  
of Technology

47

## Shadow Map



Depth buffer (or shadow map) from  
the light's point-of-view

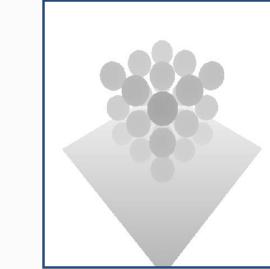


Figure Source: Justin Reynolds

Georgia Institute  
of Technology

48

## Shadow Map Concept

Compare  $z'$  with  $z_M = \text{shadow\_map}(x', y')$   
(the  $z'$  is deeper than  $z_M$  in this example)

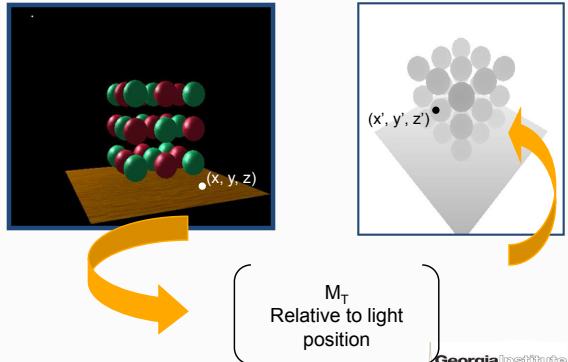


Figure Source: Justin Reynolds

$M_T$   
Relative to light  
position

Georgia Institute  
of Technology

49

## Shadow Map

Depth testing from the light's position

Two pass algorithms

1. Render shadow map (depth buffer) from the light's point of view to indicate the depth of the closest pixels to the light
  - Using specialized vertex and pixel shaders
  - Pixel shader writes pixel depth (not color)
2. Render scene from viewer's position, for each rasterized fragment
  - Determine fragment's XYZ position relative to the light
  - Compare the fragment's depth and the depth stored in the shadow map
  - The pixel is in shadow if  $P_z > \text{shadow\_map}(p_x, p_y)$
  - Can only show ambient+emissive light color for shadowed pixel

Georgia Institute  
of Technology

50

## Shadow Map Algorithm

```
procedure SHADOWMAPPING
  Render depth buffer (Z-buffer) from lights point of view,
  resulting in a shadow map or depth map
  Now, render scene from the eye's point of view

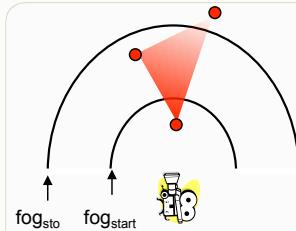
  for all rasterized fragments do
    Determine fragment's xyz position relative to the light
    That is transform each fragment's xyz into the light's coordinate
    system
    A = depth map(x, y)
    B = z-value of fragment's xyz light position
    if A < B then
      fragment is shadowed
    else
      fragment is lit
    end if
  end for
```

Algorithm Source: Justin Reynolds

Georgia Institute  
of Technology

51

## Apply Fog Effect



$$\text{color} = (1-f) \cdot \text{Color}_{\text{vertex}} + f \cdot \text{Color}_{\text{fog}}$$

$$f = \text{MAX}\left(\frac{\text{dist}(\text{eye}, \text{vertex}) - \text{fog}_{\text{start}}}{\text{fog}_{\text{stop}} - \text{fog}_{\text{start}}}, 0\right)$$

- Calculate distance
- Calculate intensity of vertex color based on distance
  - Color blending
  - Linear density, exponential density
- Blending color schemes
  - Per-vertex (then interpolate pixels), less expensive
  - Per-fragment basis (Nvidia hardware), better quality
- Color blending

Georgia Institute  
of Technology

52